

CS445/545: Homework #1

Due on October 25, 2020 at 3:10pm

Dr. A. Rhodes

Austen Nelson

Problem 1

Experiment 1: Vary number of hidden units

This scatter plot shows how changing the number of neurons in the hidden layer affected the rate of training of the network. The dotted lines represent the training set and the solid lines are the corresponding test sets. Values of 20, 50, and 100 were chosen for the dimension of the hidden layer. The learning rate was fixed to 0.1, momentum to 0.9, and weights were initialized to random values between -0.05 and 0.05. The networks were trained for 50 epochs and the accuracy as well as a confusion matrix was recorded at the end of each epoch on both a disjoint test set and the training set. The confusion matrices shown are from epoch 10, as the networks didn't improve after that point.




Figure 1: hidden scatter




Figure 2: hidden confusion

- 1. How does the number of hidden units affect the final accuracy on the test data?** For the three levels of hidden units I tested, increasing the number of hidden units correlated with an increase in overall accuracy of the network. This is expected because it allows for more complex neural pathways to emerge during training. Looking at the confusion matrices, the 20 and 50 unit networks had a difficult time identifying 5's by mistaking them as 3's. The 100 unit network seems to have figured this difference out but at the cost of minor loss of accuracy of other numbers. It is also interesting to note that 20 and 50 would mistake 5's as 3's but rarely 3's as 5's, where the 100 unit network had similar accuracy between these pairs.
- 2. How does it affect the number of epochs needed for training to converge?** It seems that increasing the number of hidden units slightly increased the amount of training needed to converge the network. 20 hit near its best on the 5th epoch, 50 on the 8th, and 100 on the 10th. 20 and 50 both had accuracies that barely exceeded these early maximums in later epochs but after significant oscillation in accuracy. This leads me to believe that these later peaks aren't meaningful. The oscillation in accuracy after the initial maximums decreases as number of hidden neurons increased, but none of the networks really benefited from more than 10 epochs of training on the test set.

3. Is there evidence that any of your networks has overfit to the training data? If so, what is that evidence? There is evidence that all 3 of the networks were overfit. Each network did not improve much after 10 epochs on the test set, and the 20 got worse. For all the networks the accuracy on the training set gradually increased well beyond the 10th epoch demonstrating some level of memorization happening.

4. How do your results compare to the results obtained by your perceptron in HW 1? Which homework?

Problem 2

Experiment 2: Vary the momentum value

This scatter plot shows how changing the momentum value affected the rate of training of the network. Values of 0, 0.25, 0.5, and 0.9 were chosen for the momentum coefficient for each network. Hidden units was fixed at 100, the learning rate at 0.1, and weights were initialized to values between -0.05 and 0.05. The confusion matrices were taken from the 10th epoch.




Figure 3: momentum scatter

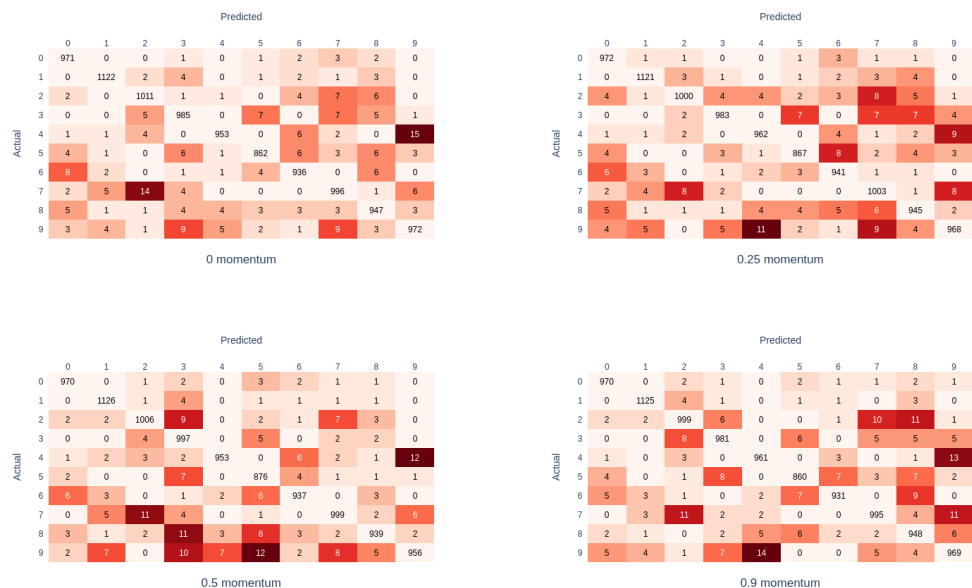


Figure 4: momentum confusion

How does the momentum value affect the final accuracy on the test data? The momentum value had little effect on the maximum accuracy of the network, but higher momentum was more prone to

overfitting. The 0.9 momentum started to diverge from the other networks in accuracy around 10 epochs. This seems intuitive as a higher momentum is going to begin memorizing the training set faster.

How does it affect the number of epochs needed for training to converge? For the first 3 to 5 epochs, the larger momentums approached their maximum accuracy faster. This is also intuitive. Between 5 and 10 epochs, though, there seems to be little difference.

Again, is there evidence that any of your networks has overfit to the training data? Again, there is little improvement after 10 epochs. This suggests overfitting because the accuracy of the training set continues to improve where the accuracy of the test sets oscillates or decreases. As mentioned in the first question, it seems that higher momentums are more prone to overfitting.

Problem 3

Experiment 3: Vary the number of training examples

This scatter plot shows how changing the number of training examples affects the rate of training of the network. $\frac{1}{4}$ (15,000), $\frac{1}{2}$ (30,000), and all (60,000) of the examples were used for each network. Hidden units was fixed at 100, momentum to 0.9, and weights were initialized to values between -0.05 and 0.05. The confusion matrices of the training set were checked to make sure an even proportion of examples was used in each network. The confusion matrices shown are from the 10th epoch of the test set on these networks.




Figure 5: test set size scatter




Figure 6: test set size confusion

How does the size of the training data affect the final accuracy on the test data? Reducing the number of training examples has a clear affect of reducing the maximum accuracy of the network. Although the half and full set are fairly close, there is nearly an entire percent difference in accuracy from the quarter network.

How does it affect the number of epochs needed for training to converge? The networks with more training examples also converged at the maximum accuracy faster. The full train set network was

nearing its maximum as early as the 4th epoch, the half around the 6th, and the quarter peaked at 12. This is also intuitive because with more data you are doing more adjustment per epoch.

Again, is there evidence that any of your networks has overfit to the training data? Again, beyond 10 all the networks started to do worse on test accuracy but training accuracy slowly increased. I am surprised the training accuracy for the quarter network took the longest to level off because there is less data to memorize but it makes sense that less training data is simply just less adjustment per epoch. The quarter network was significantly more susceptible to overtraining and really started to fall off earlier than the other two networks.